

Image Classification

fine-Tuning on resnet18

Environment

- 程式碼位於[Google Colab](#)
- Google Colab是什麼?
 - 以Jupyter notebook呈現
 - 操作位於雲端的虛擬機
 - 提供基本環境 (Torch, numpy等等)
 - 可以共用分享



 PyTorch

Google Colaboratory



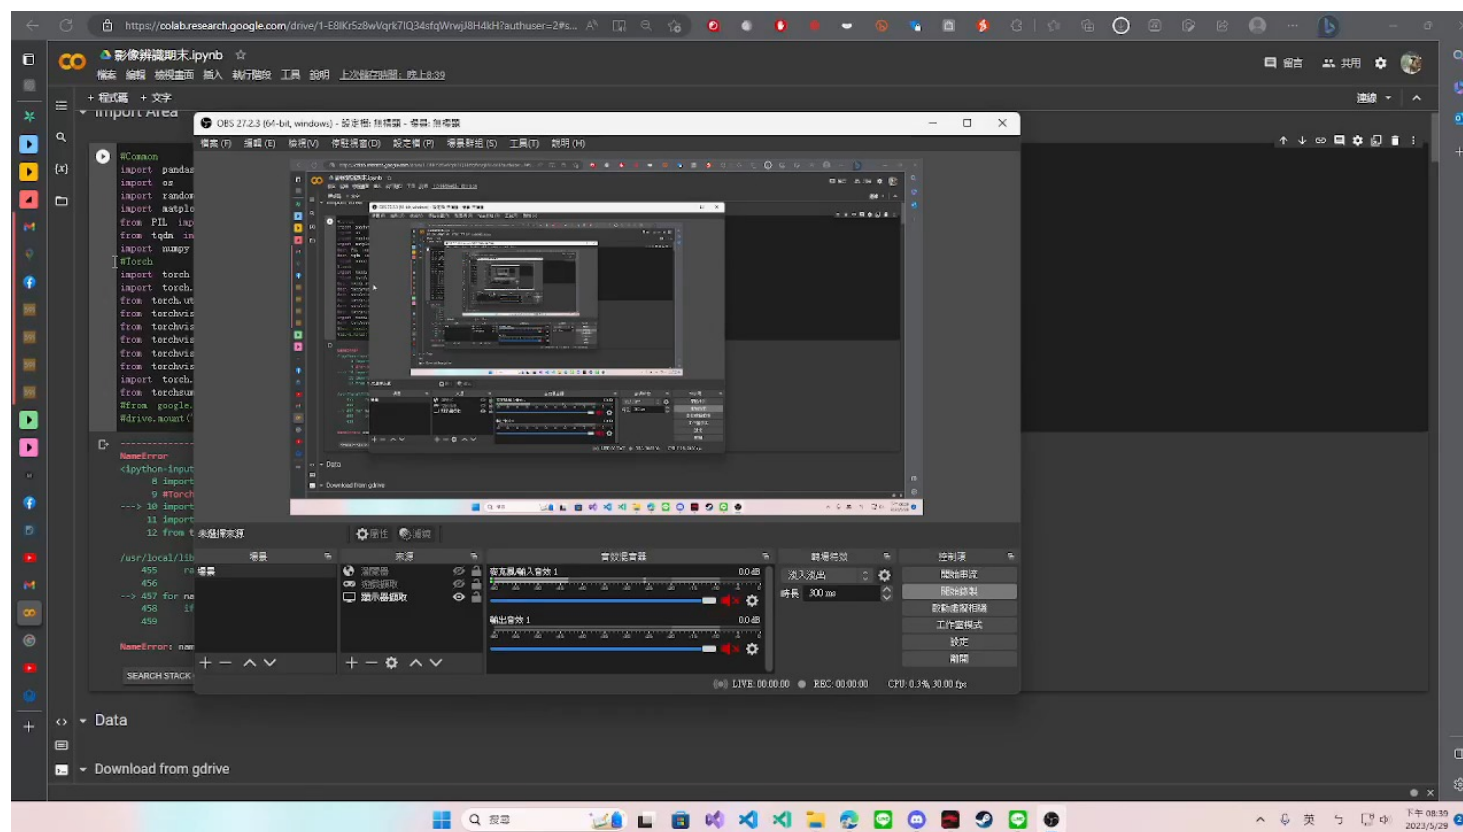
Dataset

- Link on Kaggle: [Dataset](#)
- 內有53張撲克牌 (各個花色的A~13 + Joker)
 - 每一張牌都有120~160張圖片
- 共有約8000張圖片
 - Train: 7624
 - Test: 265
 - Validation: 265



Demo

- [Link](#)



Import

```
[ ] #Common
import pandas as pd
import os
import random
import matplotlib.pyplot as plt
from PIL import Image
from tqdm import tqdm
import numpy as np
#Torch
import torch
import torch.nn as nn
from torch.utils.data import Dataset, DataLoader
from torchvision import datasets
from torchvision.models import resnet18
from torchvision.transforms import ToTensor
from torchvision.io import read_image
from torchvision import transforms
import torch.optim as optim
from torchsummary import summary
#from google.colab import drive
#drive.mount('/content/gdrive')
```

Custom Dataset

```
class CustomImageDataset(Dataset):  
    def __init__(self, dataframe, img_dir, transform=None, target_transform=None):  
        self.img_labels = dataframe  
        self.img_dir = img_dir  
        self.transform = transform  
        self.target_transform = target_transform  
  
    def __len__(self):  
        return len(self.img_labels)  
  
    def __getitem__(self, idx):  
        #Image: Path -> PIL.Image -> Tensor(Tranfrom)  
        #Label: int -> Tensor(Transform)  
        img_path = os.path.join(self.img_dir, self.img_labels.iloc[idx, 1])  
        image = Image.open(img_path)  
        label = self.img_labels.iloc[idx, 0]  
        if self.transform:  
            image = self.transform(image)  
        if self.target_transform:  
            label = self.target_transform(label)  
        return image, label
```

Data Preprocess

- 以下是有使用到的前處理Method

- Resize

- 隨機水平翻轉

- 隨機垂直翻轉

- ToTensor

- Normalize


```
tform_train = transforms.Compose([
    transforms.Resize(224),
    transforms.RandomHorizontalFlip(p=0.5),
    transforms.RandomVerticalFlip(p=0.5),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
])

tform_test = transforms.Compose([
    transforms.Resize(224),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
])
```

Model


- Detail
 - 名稱：ResNet18
 - Pre-trained on ImageNet

```
model = resnet18(weights='DEFAULT')
model.fc = nn.Sequential(
    nn.Dropout(p=0.5),
    nn.Linear(512, len(class_labels))
)
model
```



Residual Network

- Residual (殘差) 的好處
 - 避免梯度消失
 - 可以增加深度
 - 模型表現變好
 - 選擇較好的Layer



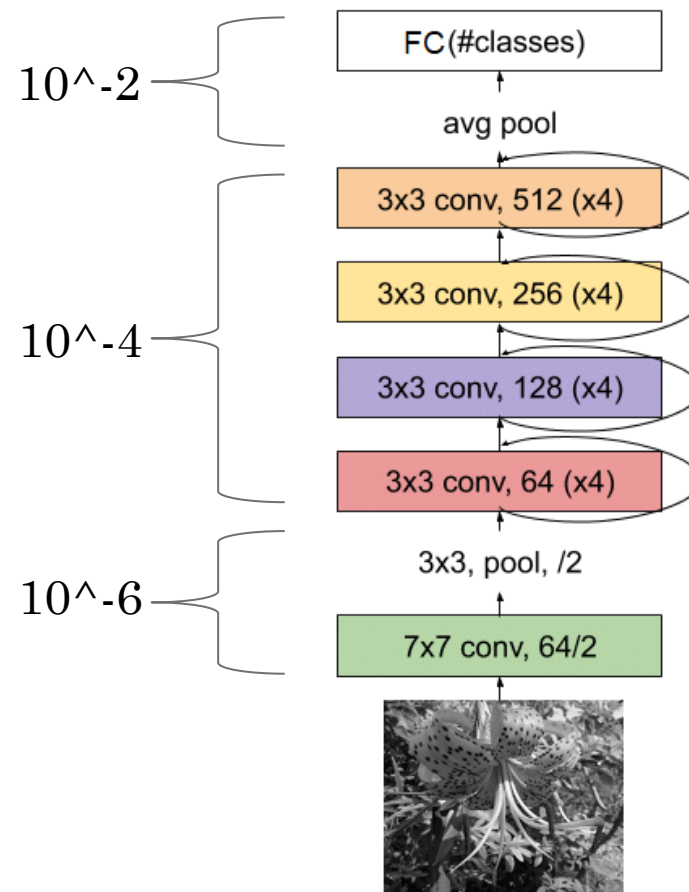
Hyper Parameter

- Epoch: 30
- Batch size: 32
- Optimizer: Adam
 - Learning rate: 10^{-6} , 10^{-4} , 10^{-2}
 - Scheduler step: 30
- Early stop: 3

```
class Config:
    epoch = 30
    batch_size = 32
    num_workers = 0
    lr = [1e-6, 1e-4, 1e-2]
    lr_step = 30
    device = 'cuda' if torch.cuda.is_available() else 'cpu'
    patience = 3
    checkpoints = "./4.pth"
    test_model = ""
```

Fine Tuning

- 樣本數較少時使用
- 方法
 - 凍結Layer
 - 給予每個Layer不同的Learning rate
 - 底層 -> 10^{-6}
 - 中間的卷積層 -> 10^{-4}
 - 分類層 -> 10^{-2}



Data Loader(Code)

```
print(len(train_ds))  
train_dataLoader = DataLoader(train_ds, batch_size=Config.batch_size, shuffle=True,  
                               num_workers=Config.num_workers, drop_last=True)
```

7624

```
print(len(valid_ds))  
valid_dataLoader = DataLoader(valid_ds, batch_size=Config.batch_size, shuffle=True,  
                               num_workers=Config.num_workers, drop_last=True)
```

265

Checkpoint(Code)

```
▶ need_to_train = False #Set true if you want to continue training
last_epoch = -1
try:
    # Using when training
    # weight_list = os.listdir(Config.checkpoints)
    # weight_list = [w.split('.')[0] for w in weight_list]
    # lastest_weight = f"{weight_list[-1]}.pth"
    # best_weight_dir = os.path.join(Config.checkpoints, lastest_weight)

    # Load checkpoint(weight, last_epoch, loss) in to model
    best_weight_dir = "./4.pth"
    checkpoint = torch.load(best_weight_dir, map_location=Config.device)
    model.load_state_dict(checkpoint['model_state_dict'])
    last_epoch = checkpoint['epoch']
    loss = checkpoint['loss']
    Config.test_model = best_weight_dir
    print(f"Load best weight path:{best_weight_dir}, epoch:{last_epoch} / {Config.epoch}, loss:{loss}")
except Exception as e:
    print(e)
    need_to_train = True
```

Fine Tuning(Code)

```
#Param1 -> bottom few layers 10^-6
#Param2 -> others 10^-4
#Param3 -> addition layer 10^-2
modelParams = []
temps = []

for param in list(model.parameters())[0:3]:
    temps.append(param)

modelParams.append(temps)

temps = []
for param in list(model.parameters())[3:-2]:
    temps.append(param)

modelParams.append(temps)

temps = []
for param in list(model.parameters())[-2:]:
    temps.append(param)

modelParams.append(temps)
```

Optimizer(Adam) Code

```
▶ # Training Setup

# Loss -> cross entropy
criterion = nn.CrossEntropyLoss()

# Optimizer -> Adam
optimizer = optim.Adam(
    [
        {'params': modelParams[0], 'lr':Config.lr[0]},
        {'params': modelParams[1], 'lr':Config.lr[1]},
        {'params': modelParams[2], 'lr':Config.lr[2]},
    ],
    lr=Config.lr[1],
)
scheduler = optim.lr_scheduler.StepLR(optimizer, step_size=Config.lr_step, gamma=0.1)
```

Train Code (Train dataset)

```
if need_to_train:
    model.to(Config.device)
    best_loss = 100 #For valid set
    count = 0

    # Iterate
    for e in range(last_epoch + 1, Config.epoch):
        # Train
        correct_count = 0
        train_last_epoch = e
        model.train()

        with tqdm(train_dataLoader, desc=f"Epoch {e}/{Config.epoch}", total=len(train_dataLoader)) as progress_bar:
            for data, labels in progress_bar:
                #To cpu/gpu
                data = data.to(Config.device)
                labels = labels.to(Config.device)
                #Gradient descent
                optimizer.zero_grad()
                #Output
                outputs = model(data)
                #Gradient descent by loss
                train_loss = criterion(outputs, labels)
                train_loss.backward()
                #Optimizer
                optimizer.step()

                #Get prediction and count accuracy
                _, preds = torch.max(outputs, 1)
                correct_count += torch.sum(preds == labels.data)

            progress_bar.set_postfix({"loss": train_loss.item()})

        acc = correct_count.double() / (len(train_dataLoader) * Config.batch_size)

        print(f"Epoch {e}/{Config.epoch}, Train loss: {train_loss}, acc: {acc}")
        #Record
        training_loss.append(train_loss)
        training_acc.append(acc)
```


Train Code (Valid dataset)

```
#Evaluation mode
correct_count = 0
torch.cuda.empty_cache()
model.eval()

with tqdm(valid_dataLoader, desc=f"Epoch {e}/{Config.epoch}", total=len(valid_dataLoader)) as progress_bar:
    for data, labels in progress_bar:
        # To cpu/gpu
        data = data.to(Config.device)
        labels = labels.to(Config.device)
        # Get output
        outputs = model(data)
        # Loss cal
        val_loss = criterion(outputs, labels)

        # Predicitons
        _, preds = torch.max(outputs, 1)
        correct_count += torch.sum(preds == labels.data)

    progress_bar.set_postfix({"loss": val_loss.item()})

acc = correct_count.double() / (len(valid_dataLoader) * Config.batch_size)
print(f"Epoch {e}/{Config.epoch}, Validation loss: {val_loss}, acc: {acc}")
# Record
validation_loss.append(val_loss)
validation_acc.append(acc)
```

Train Code (Save weight)

```
# If loss smaller than before
if val_loss < best_loss:
    count = 0
    best_loss = val_loss
    # At this point also save a snapshot of the current model
    backbone_path = os.path.join(Config.checkpoints, f"{e}.pth")
    torch.save({
        'epoch': e,
        'model_state_dict': model.state_dict(),
        'optimizer_state_dict': optimizer.state_dict(),
        'loss': val_loss,
    }, backbone_path)
    Config.test_model = backbone_path
else:
    # Early stop
    count += 1
    if count >= Config.patience:
        break


#Clean cache
del train_loss, val_loss, outputs
torch.cuda.empty_cache()

scheduler.step()
```

Loss & Accuracy



Early Stop



註 : Training with dropout($p=0.5$)

Test dataset(Code)

```
▶ correct_count = 0
model.eval()
model.to(Config.device)

with tqdm(test_dataLoader, total=len(test_dataLoader)) as progress_bar:
    for data, lbl in progress_bar:
        #To cpu/gpu
        data = data.to(Config.device)
        lbl = lbl.to(Config.device)
        #Get output
        outputs = model(data)
        #Cal loss
        loss = criterion(outputs, lbl)
        #Get prediction
        _, preds = torch.max(outputs, 1)
        correct_count += torch.sum(preds == lbl.data)

print(f"Loss: {loss}, acc: {correct_count.double() / (len(test_dataLoader))}")
```

100%|██████████| 265/265 [00:25<00:00, 10.25it/s]Loss: 0.0010851691477000713, acc: 0.9320754716981132

Show Result 3x3(Code)

```
▶ w = 10
  h = 10

columns = 3
rows = 3

fig = plt.figure(figsize=(8, 8))
ax = []

for i in range(1, columns*rows + 1):
    #Random fetch image in test dataset
    idx = random.randint(0, len(test_ds)-1)
    show_img, show_lbl = test_ds[idx]
    show_img, show_lbl = show_img.to(Config.device), show_lbl

    #Fit input shape
    x = show_img
    x = x.unsqueeze(0)

    #Evaluation
    model.eval()
    model = model.to(Config.device)

    #Get output without gradient descend
    with torch.no_grad():
        out = model(x)

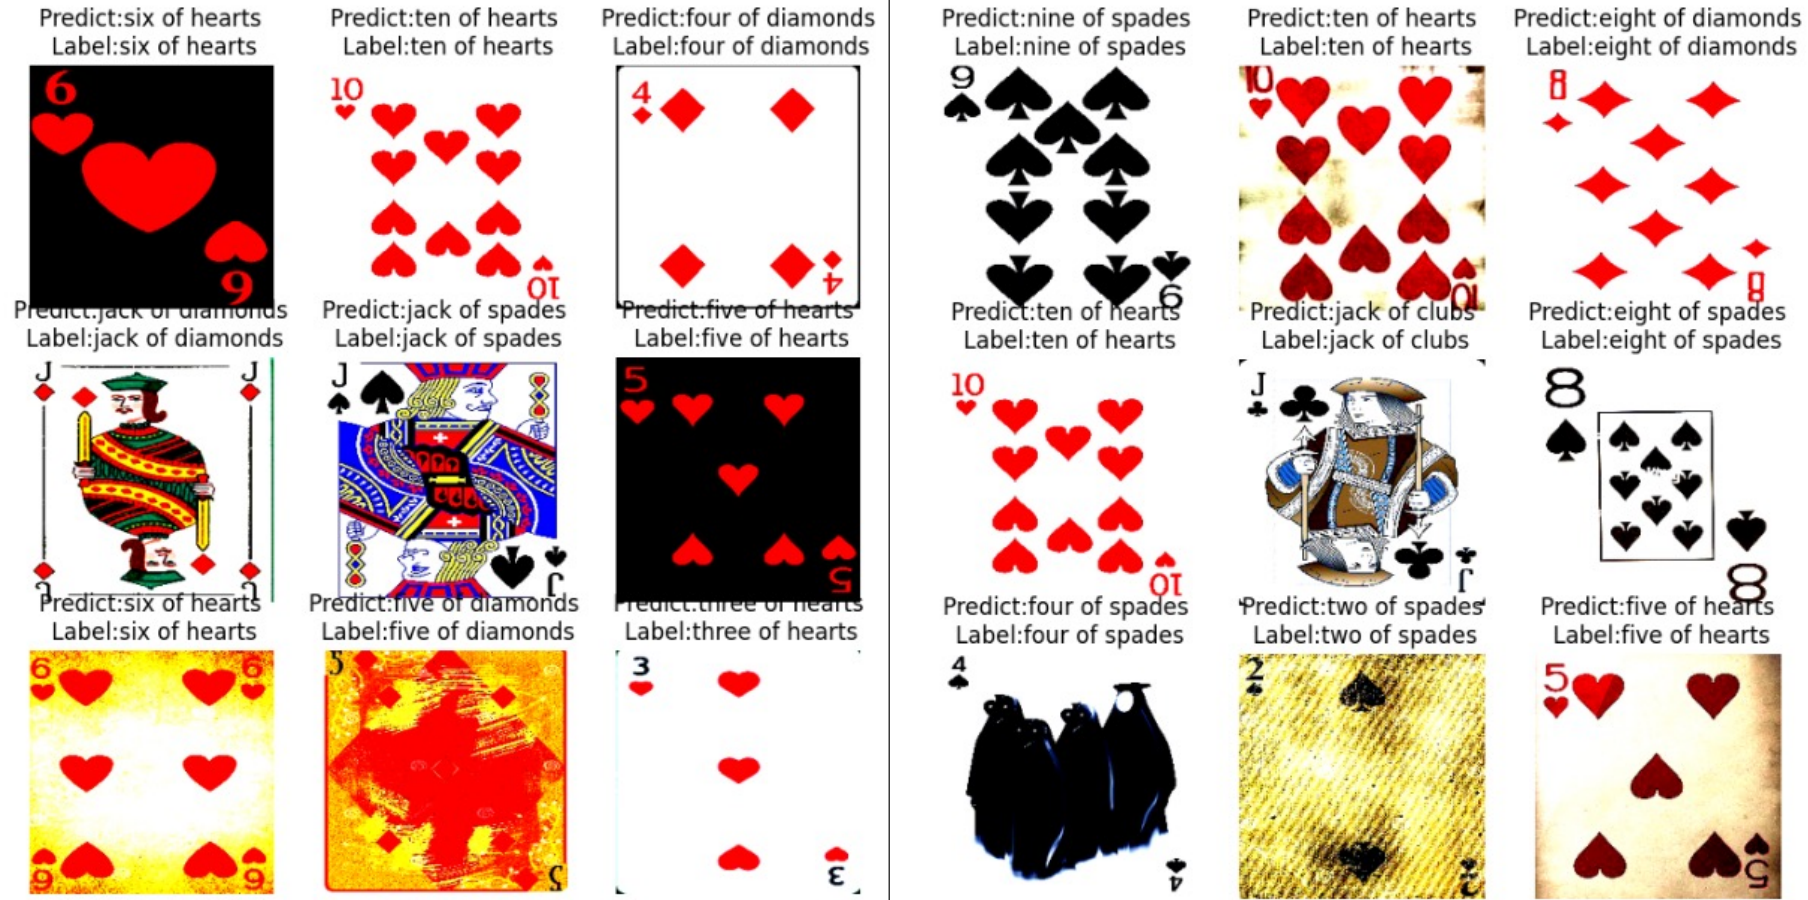
    show_img, show_lbl, out = show_img.cpu(), show_lbl, out.cpu()

    #Show in subplot
    ax.append(fig.add_subplot(rows, columns, i))
    ax[-1].set_title(f"Predict:{class_labels[np.argmax(out)]}\n Label:{class_labels[show_lbl]}")
    plt.imshow(show_img.permute(1, 2, 0))

plt.axis('off')

plt.show()
```

Show Result(3x3)



Try on other Image


Predict:two of diamonds



Predict:ace of hearts



Predict:ace of spades



Special Case(Different style)

Predict:six of spades



Predict:seven of hearts




Predict:three of spades



Special Case(Different angle)

Predict:seven of spades



Predict:jack of spades



Predict:ace of hearts



Predict:jack of spades



Special Case(Different angle)

Predict:ace of spades



Predict:four of spades



Predict:four of clubs



Special Case(For Fun)

Predict:three of spades



Predict:ace of spades



Predict:four of clubs



Improvement

- On dataset
 - 裁切重點部分
 - 資料篩選
- On model
 - 結合其他模型
 - Yolo
 - OCR
 - 增加模型深度

Application

- AI牌類遊戲（真實世界）
- 撲克牌品質檢定
- 撲克牌整理
- 發牌機器（及相關統計）

Reference

- [Residual network](#)
- [Pytorch document](#)
- Transfer learning on ResNet50 (有印象但找不到)

Q&A

Thank you
for
listening!



Author: LyraEri